

LANGUAGE DRIVEN OUTFIT STYLE ENHANCEMENT

...

Team 09

Shrey Shrivastava(17CS30034)

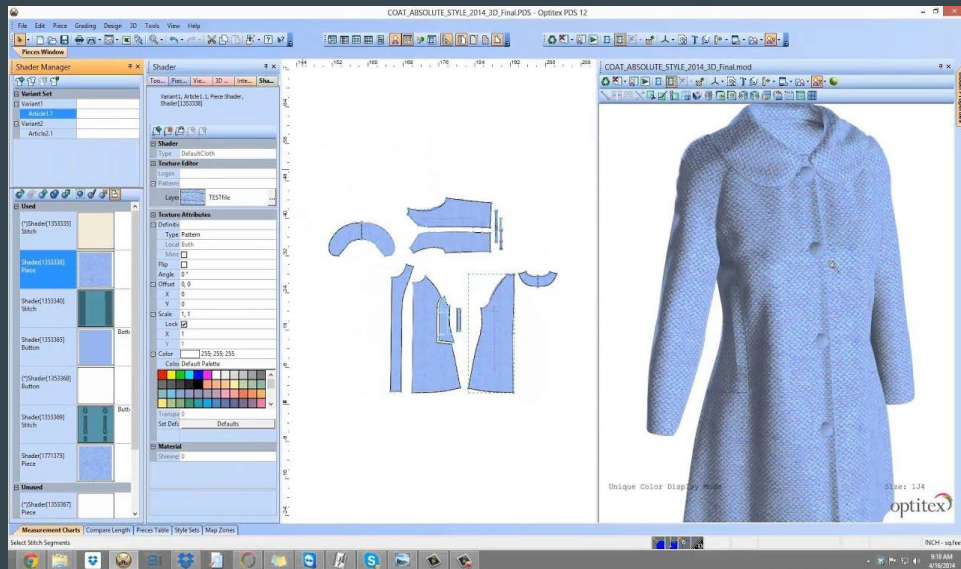
Snehal Reddy Koukuntla(17CS30020)

Deepank Agrawal(17CS30011)

Kanishk Singh(17CS30018)

Motivation

- A big challenge in the fashion and graphics industry is for any person to stylise their self-designed clothes.
- A natural way for a layman to describe his idea is through words and hence the best way to take input is natural language in form of textual descriptions.



Formal Statement and Plan

- Let's formally present the problem statement we seek to solve:
 - Given an input image of a fashion item I , and a textual description consisting of a set of key-words, we aim to generate transformed image T , with the style corresponding to the textual description imposed on top of the input image I .
- We divide the problem statement into **three** major parts to ease task formulation:
 - Generate the texture from the textual description input from the user.
 - Segmentation of targeted cloth item from input image.
 - Generated texture is superimposed on the segmented region.

Step 1: Texture generation from input textual description

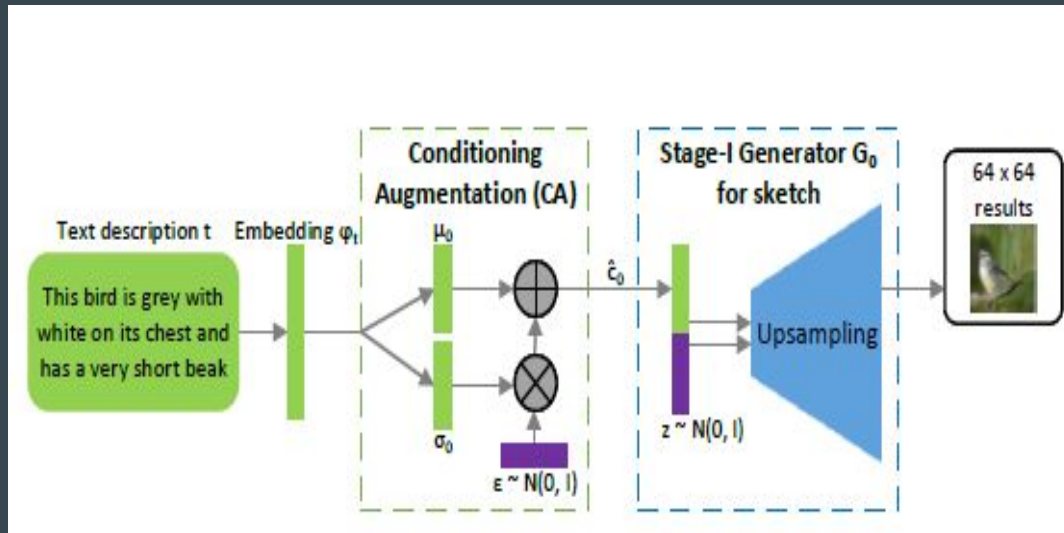
- We employ StackGAN for the task of texture generation.
- GANs have performed effectively in generating images from textual descriptions, however the synthesized images lack details and vivid object parts.
- To tackle the challenges , the proposed approach decomposes the problem into two more manageable sub-problems.
 - A low resolution image is generated using STAGE-I GAN conditioned on text description.
 - STAGE-II GAN generates realistic high resolution images conditioned on the low resolution image and the corresponding text description.
- Since we only need a patch for our original task , we only deployed the STAGE-I GAN in our algorithm.
- The image output from the first stage is a low resolution image consisting of rough shape and basic colours conditioned on textual description.

Stage-I Generative Adversarial Networks (GANs)

- Composed of two models that are alternatively trained to compete with each other
 - The Generator G
 - Optimized to generate images that are difficult for the discriminator D to differentiate from real images.
 - The Discriminator D
 - Optimized to distinguish real images from the synthetic images generated by G .


Stage-I Generator

- Text description \mathbf{t} is encoded to form text embedding.
- Conditional Augmentation is added to improve the diversity of generated images.
- To avoid over-fitting, **KL divergence** between standard and conditioning Gaussian distribution is added during the generator training.




Stage-I Discriminator

- Text embedding is initially compressed and spatially replicated to form tensor of shape $[4,4,128]$.
- The image is also down sampled from $[64, 64, 3]$ to $[4, 4, 1024]$.
- The resulting concatenated tensor is later fed to $[1,1]$ convolution layer followed by a $[4,4]$ convolution layer.
- The decision of this discriminator is done by a fully connected layer with one node.




Results

- Sample images generated from STACKGAN



“Wooden Door”




“Furniture”



“Cheese Pizza”



“Christmas Tree”



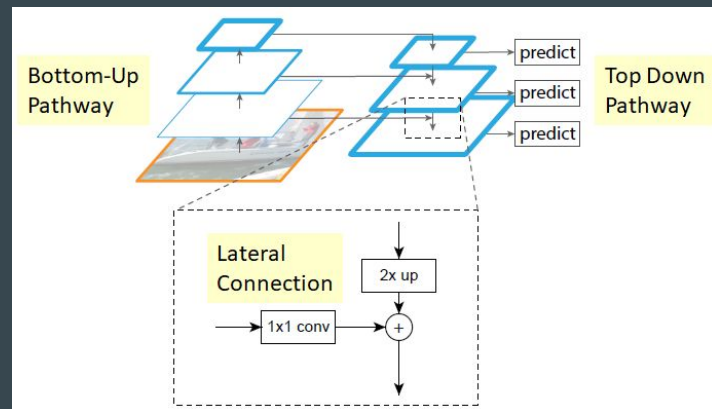
“Beach”

Step 2: Segmentation of cloth piece

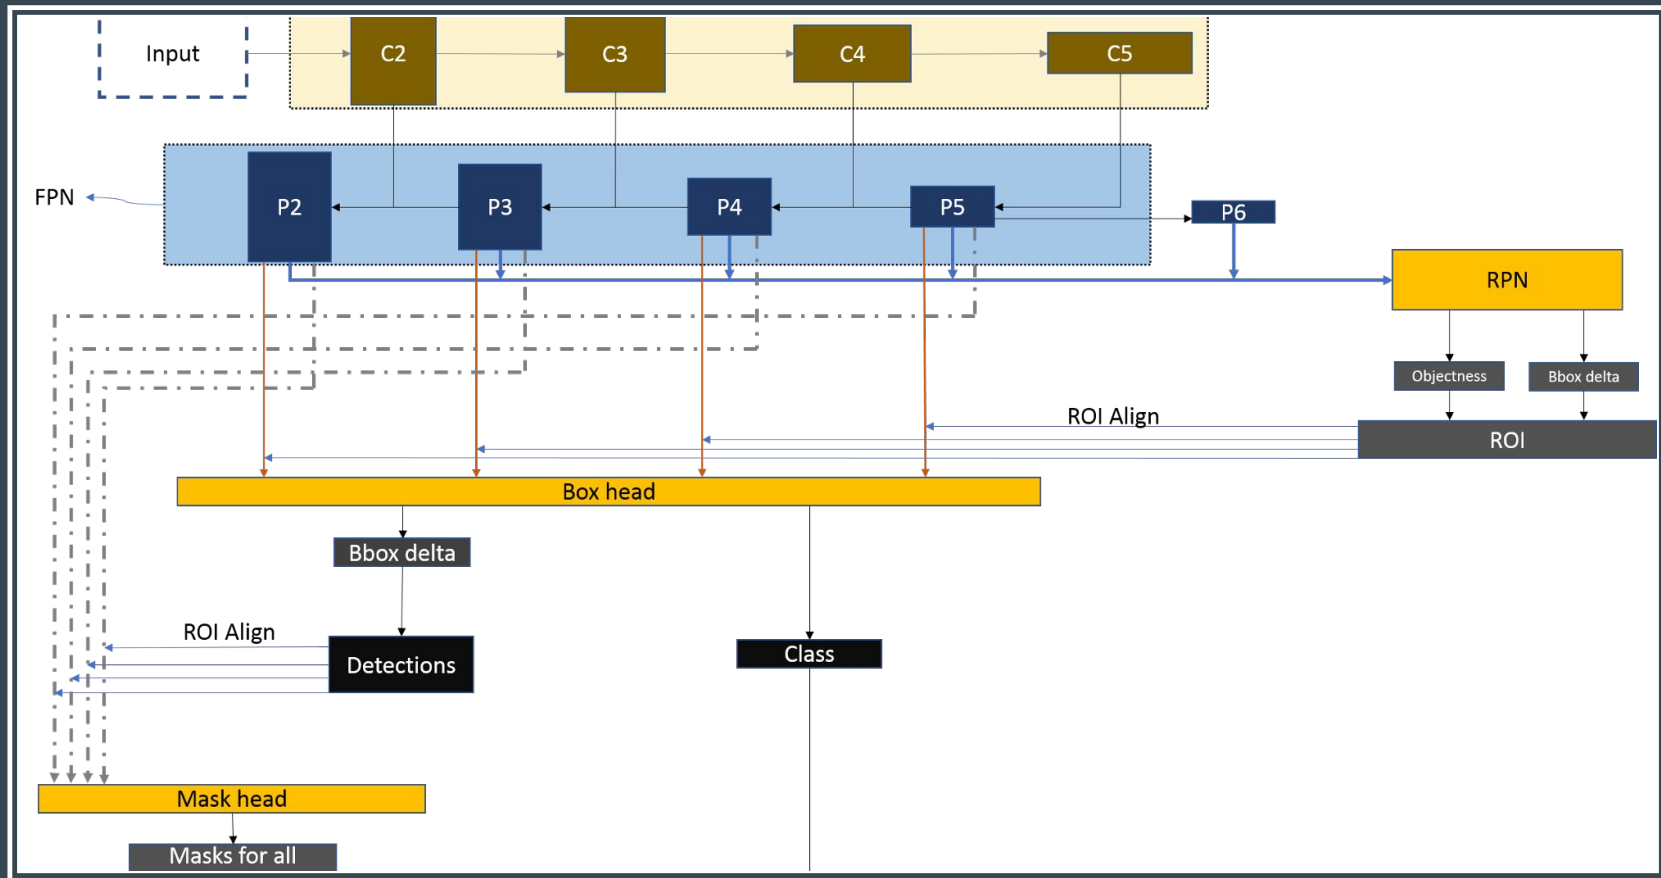
- We employ **Mask R-CNN** for the task of image segmentation.
- Mask R-CNN was initialised with **ResNet-50** backbone pre-trained model for segmentation on MS COCO dataset.
- **DeepFashion 2** is a comprehensive fashion dataset. It contains 491K diverse images of 13 popular clothing categories from both commercial shopping stores and consumers.
- We used this dataset to further fine-tune the Mask RCNN. This can segment different classes of items separately on the same image like the top, bag, boots, pants etc.
- The resultant segmentation is then passed on to the next part of the pipeline.

Architecture


- Mask RCNN has three major differences:
 - Use of FPN's
 - Replacement of ROI-Pool with ROI-Align
 - Introduction of additional branch to generate masks.
- The extra mask head allows us to pixel wise segment each object and also extract each object separately without any background
- The Mask RCNN architecture we chose had a Resnet 50 backbone with Feature Pyramid Network.



```
def get_model_instance_segmentation(num_classes):  
    # load an instance segmentation model pre-trained pre-trained on COCO  
    model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)  
  
    # get number of input features for the classifier  
    in_features = model.roi_heads.box_predictor.cls_score.in_features  
    # replace the pre-trained head with a new one  
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)  
  
    # now get the number of input features for the mask classifier  
    in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels  
    hidden_layer = 256  
    # and replace the mask predictor with a new one  
    model.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask,  
                                                         hidden_layer,  
                                                         num_classes)  
  
    return model
```



Example: Segmentation



Step 3: Texture application over the segmented image

- After the segmentation of the target cloth piece, we superimpose the texture over the segmented region.
- The size of resulting image is same as the input image.
- The resulting image is passed into a pre-trained **TextureGAN** model.

Step 3: Continued...

- Since a generalised pre-trained model was available we decided not to train from scratch. We used 'clothe' pre-trained model of TextureGAN.
- The pretrained model is trained on DeepFashion Dataset.
- The generated image is combined with the input image to generate final output.

TextureGAN

- First deep image synthesis method which allows user to implicitly control object texture.
- The network realistically applies these textures to the target object.
- The model is trained in two stages:
 - Ground-truth Pre-training.
 - External texture Fine-tuning.

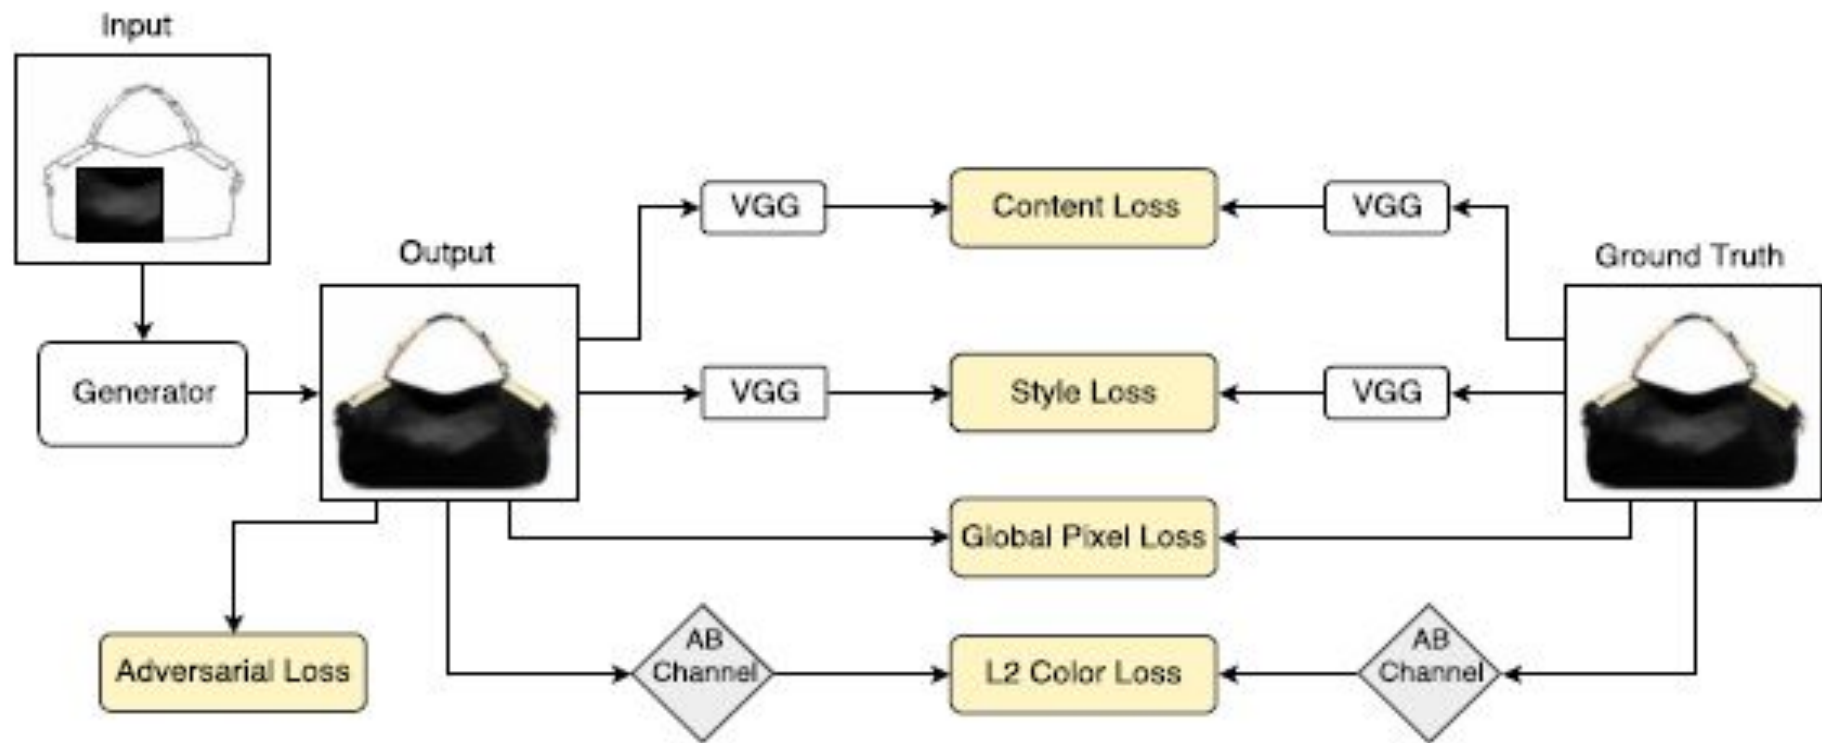


Published results from the paper

Ground-truth Pre-training

- The aim of this stage is propagate the texture information in the entire object.
- The input image is converted into LAB color space.
- Various losses are introduced to ensure controlled texture generation:
 - Losses imposed on L channel:
 - Feature loss: To control “leaking” of texture from boundaries.
 - Adversarial loss: Standard GANs loss.
 - Texture loss: To control the variation in generated textures.
 - Style loss: To propagate texture details as per the input texture patch.
 - Pixel loss: Claimed to stabilise the training and generation of faithful textures.
 - Loss imposed on AB channels:
 - Color loss: To enforce user’s color constraints.

$$\mathcal{L} = \mathcal{L}_F + w_{ADV} \mathcal{L}_{ADV} + w_S \mathcal{L}_S + w_P \mathcal{L}_P + w_C \mathcal{L}_C$$



TextureGAN pipeline for the ground-truth pre-training

External Texture Fine-tuning

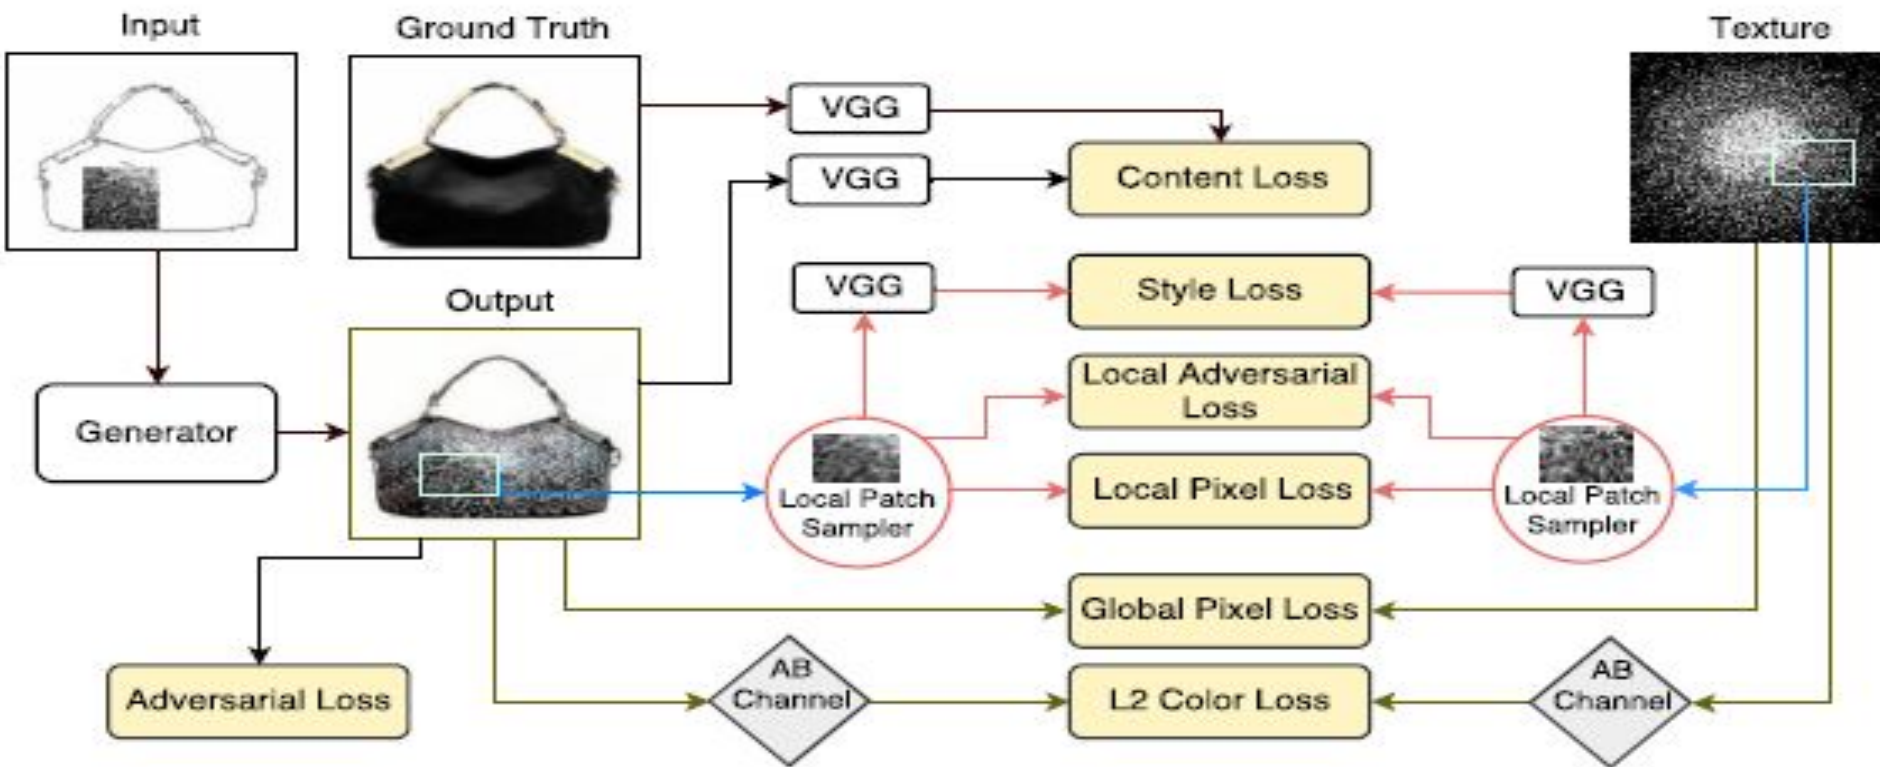
- The network is trained on limited textures in the 1st step. And it has been observed that network performs poorly in propagating rare high level texture details.
- In this step, TextureGAN is generalised to support a broader range of textures and to propagate unseen textures better by fine-tuning the network with a separate texture-only database.
- The main difference is that, in the 1st step, model was “faithful” to the output image while in this step the model is “faithful” to the input textures.

Re-definition of Loss functions

- To ensure “faithfulness” towards the input texture, some loss functions has to be re-defined.
- Feature loss and Adversarial loss remains unchanged.
- Pixel loss and Color loss are changed to compute loss between the texture and generated image.
- For better propagation of texture, Local texture loss is introduced:
 - Local Adversarial loss.
 - Local Style loss.
 - Local Pixel loss.

$$\mathcal{L}_t = \mathcal{L}_s + w_p \mathcal{L}_p + w_{adv} \mathcal{L}_{adv}$$

$$\mathcal{L} = \mathcal{L}_F + w_{ADV} \mathcal{L}_{ADV} + w_P \mathcal{L}'_P + w_C \mathcal{L}'_C + \mathcal{L}_t$$



TextureGAN pipeline for the External Texture Fine-Tuning

Results


Image 1



Input image



The StackGAN
output with
input “Purple
Flower”.



The final output with
input “Purple Flower”

Conclusion

- With the presented system, we have successfully combined a wide range of recent research developments to produce a tool where a user can textually describe the details of generated stylised clothing.
- Our results show that this pipeline can handle a wide variety of textual inputs and generate texture compositions that follow the sketched contours satisfactorily.
- Still the proposed methodology is not end-to-end trainable. We have used pre-existing models and cascaded them to generate the output images.
- For future research towards an end-to-end model, our analysis of related research work suggests using extra losses and degenerator modelling in vanilla C-GAN would be necessary for achieving the desired results.

Thank You